6. Explore conditional probabilities, independence, and expected outliers. Examine distribution shapes, skewness, correlations, and compare different outlier detection methods to assess consistency and insights.

✓ Clean & Compact R Code

```
# Load libraries
library(dplyr)
library(ggplot2)
library(e1071)
library(car)
# Load dataset
data <- mtcars
data$cyl <- as.factor(data$cyl)</pre>
data$gear <- as.factor(data$gear)</pre>
# 1. Conditional Probability: P(am = 1 | gear = 4)
cond prob <- mean(data$am[data$gear == 4])
cat("P(am = 1 | gear = 4):", round(cond prob, 3), "\n")
# 2. Independence: Chi-squared test between am and gear
chi p <- chisq.test(table(data$am, data$gear))$p.value
cat("Chi-squared p-value (am vs gear):", round(chi p, 4),
  ifelse(chi p < 0.05, "\rightarrow Dependent\n", "\rightarrow Independent\n"))
#3. Skewness
cat("Skewness of mpg:", round(skewness(data$mpg), 3), "\n")
cat("Skewness of hp:", round(skewness(data$hp), 3), "\n")
#4. Correlation matrix
cat("\nCorrelation matrix:\n")
print(cor(select if(data, is.numeric)))
```

```
# 5. Distribution of mpg
ggplot(data, aes(mpg)) +
 geom histogram(binwidth = 2, fill = "skyblue", color = "black") +
 ggtitle("MPG Distribution") + theme minimal()
#6. Outlier Detection
## (a) IQR Method
Q1 <- quantile(data$mpg, 0.25)
Q3 <- quantile(data$mpg, 0.75)
IQR val <- Q3 - Q1
cat("\nIQR Outliers (mpg):", out iqr, "\n")
## (b) Z-score Method
z scores <- scale(data$mpg)</pre>
out z \le -data pg[abs(z scores) > 2]
cat("Z-score Outliers (mpg):", out z, "\n")
## (c) Bonferroni Method
model <- lm(mpg \sim ., data = data)
cat("Bonferroni Outlier Test:\n")
print(outlierTest(model))
```

P Notes:

- scale() standardizes for Z-score.
- outlierTest() from the car package flags outliers based on studentized residuals.
- You can change mpg to other variables if needed.
- This version is **short**, **readable**, **and informative**, great for labs or reports.

7. Using Pearson and Spearman methods. Identify key influencing variables and visualize insights using heatmaps, scatter plots, box plots, bar plots, and clustering dendrograms

```
# Load necessary libraries
library(ggplot2)
library(corrplot)
library(heatmaply)
library(cluster)
library(RColorBrewer)
# Load example dataset
data(mtcars)
# 1. Pearson and Spearman Correlation
pearson corr <- cor(mtcars, method = "pearson")</pre>
spearman corr <- cor(mtcars, method = "spearman")</pre>
# 2. Pearson Correlation Heatmap
corrplot(pearson corr, method = "color", type = "upper",
     col = brewer.pal(n = 9, name = "Blues"),
     title = "Pearson Correlation Matrix")
#3. Spearman Correlation Heatmap
corrplot(spearman corr, method = "color", type = "upper",
     col = brewer.pal(n = 9, name = "RdYlBu"),
     title = "Spearman Correlation Matrix")
# 4. Scatter Plot (mpg vs wt)
ggplot(mtcars, aes(x = wt, y = mpg)) +
 geom point(aes(color = mpg)) +
 labs(title = "Scatter Plot of mpg vs wt") +
```

```
theme_minimal()
# 5. Box Plot (mpg)
ggplot(mtcars, aes(x = factor(1), y = mpg)) +
 geom boxplot() +
 labs(title = "Box Plot of mpg") +
 theme minimal()
# 6. Bar Plot (cyl)
ggplot(mtcars, aes(x = factor(cyl))) +
 geom bar(aes(fill = factor(cyl))) +
 labs(title = "Bar Plot of Cars with Different Cylinder Counts") +
 theme minimal()
#7. Clustering Dendrogram
dist matrix <- dist(mtcars)
hc <- hclust(dist matrix)</pre>
plot(hc, main = "Dendrogram of mtcars")
#8. Clustered Heatmap
heatmaply::heatmaply(mtcars, main = "Clustered Heatmap of mtcars")
8. explain how Principal Component Regression (PCR) can be used to address
multicollinearity issues while improving predictive accuracy.
# Install and load required package
if (!require(pls)) install.packages("pls", dependencies = TRUE)
library(pls)
# Load dataset
data(mtcars)
# Prepare data
```

```
X <- scale(mtcars[, -1]) # All predictors except mpg
Y <- mtcars$mpg
                        # Response variable
# Perform PCR with Cross-Validation
per model \leftarrow per(Y \sim X, \text{ validation} = \text{"CV"})
# Summary of the PCR model
summary(pcr model)
# Plot RMSEP to find optimal number of components
validationplot(pcr model, val.type = "RMSEP", main = "PCR Cross-Validation")
# Predict mpg using 3 components
predicted <- predict(pcr model, ncomp = 3)</pre>
# Compare predicted vs actual values
plot(Y, predicted, col = "blue", pch = 16,
  xlab = "Actual MPG", ylab = "Predicted MPG",
  main = "PCR: Actual vs Predicted")
abline(a = 0, b = 1, col = "red")
9. Fit linear and quadratic regression, compare their R-squared and RMSE values, and
interpret results. Identify the best-fitting model and predict sales
# Load libraries
library(ggplot2)
library(Metrics) # For RMSE calculation
# Load dataset
data(mtcars)
# Fit linear regression model
linear model \leq- lm(mpg \sim wt, data = mtcars)
```

```
# Fit quadratic regression model
quadratic model <- lm(mpg \sim wt + I(wt^2), data = mtcars)
# Summary outputs
summary(linear model)
summary(quadratic model)
# Predictions
linear pred <- predict(linear model, newdata = mtcars)</pre>
quadratic pred <- predict(quadratic model, newdata = mtcars)</pre>
# R-squared values
linear r2 <- summary(linear model)$r.squared
quadratic r2 <- summary(quadratic model)$r.squared
# RMSE values
linear rmse <- rmse(mtcars$mpg, linear pred)</pre>
quadratic rmse <- rmse(mtcars$mpg, quadratic pred)</pre>
# Print comparison
cat("Linear Model R-squared:", linear r2, "\n")
cat("Quadratic Model R-squared:", quadratic r2, "\n")
cat("Linear Model RMSE:", linear rmse, "\n")
cat("Quadratic Model RMSE:", quadratic rmse, "\n")
# Model Comparison
if (linear r2 > quadratic r2) {
 cat("Linear model provides a better fit based on R-squared.\n")
} else {
```

```
cat("Quadratic model provides a better fit based on R-squared.\n")
}
# Predict mpg when wt = 3
new data \leftarrow data.frame(wt = 3)
linear prediction <- predict(linear model, newdata = new data)
quadratic prediction <- predict(quadratic model, newdata = new data)
cat("\nPrediction from Linear Model (wt = 3):", linear prediction, "\n")
cat("Prediction from Quadratic Model (wt = 3):", quadratic prediction, "\n")
10. Compute the Pearson and Spearman correlation coefficients for a given dataset and
analyze the relationships between variables. Explain how correlation matrices assist in
feature selection and why highly correlated features might lead to redundancy in
predictive models.
# Load library
library(corrplot)
# Load dataset
data(mtcars)
# Compute Pearson and Spearman correlations
pearson corr <- cor(mtcars, method = "pearson")</pre>
spearman corr <- cor(mtcars, method = "spearman")</pre>
# Pearson heatmap
corrplot(pearson corr, method = "circle", type = "upper",
     title = "Pearson Correlation Matrix", mar = c(0, 0, 1, 0))
# Spearman heatmap
corrplot(spearman corr, method = "circle", type = "upper",
     title = "Spearman Correlation Matrix", mar = c(0, 0, 1, 0))
```

```
# Print correlation matrices
cat("Pearson Correlation Matrix:\n")
print(pearson corr)
cat("\nSpearman Correlation Matrix:\n")
print(spearman corr)
11. Visualize correlations, split data, and train a Logistic Regression model. Evaluate
using accuracy, precision, recall, and F1-score. Apply PCA for dimensionality reduction
and assess if performance improves for better
library(ggplot2); library(caret); library(corrplot); library(pls)
data(mtcars)
mtcars$mpg binary <- ifelse(mtcars$mpg > 20, 1, 0)
# Correlation plot
corrplot(cor(mtcars[, -c(1, 11)]), method = "circle")
# Train-test split
set.seed(42)
split <- createDataPartition(mtcars$mpg binary, p = 0.7, list = FALSE)
train <- mtcars[split, ]; test <- mtcars[-split, ]
# Logistic Regression
model <- glm(mpg binary ~ ., data = train, family = "binomial")
pred <- ifelse(predict(model, test, type = "response") > 0.5, 1, 0)
# Evaluation
conf <- confusionMatrix(as.factor(pred), as.factor(test$mpg binary))</pre>
cat("Accuracy:", mean(pred == test$mpg binary), "\n")
cat("Precision:", conf$byClass["Pos Pred Value"], "\n")
```

```
cat("Recall:", conf$byClass["Sensitivity"], "\n")
cat("F1 Score:", conf$byClass["F1"], "\n")
# PCA
X \leq scale(mtcars[, -c(1, 11)])
pca \le prcomp(X)
train pca <- cbind(pca$x[split, 1:2], mpg binary = train$mpg binary)
test pca <- cbind(pca$x[-split, 1:2], mpg binary = test$mpg binary)
# PCA model
model pca <- glm(mpg binary ~ ., data = train pca, family = "binomial")
pred pca <- ifelse(predict(model pca, test pca, type = "response") > 0.5, 1, 0)
conf pca <- confusionMatrix(as.factor(pred pca), as.factor(test pca$mpg binary))
cat("\nWith PCA - Accuracy:", mean(pred pca == test pca$mpg binary), "\n")
cat("F1 Score with PCA:", conf pca$byClass["F1"], "\n")
12. Formulate the problem as a multiple linear regression equation and express it in
matrix form. Using the normal equation, calculate the regression coefficients to fit the
model. Discuss how multicollinearity among predictor variables affects the reliability of
the regression model and explain how Principal Component Regression (PCR) can be
used to address multicollinearity issues while improving predictive accuracy.
library(pls)
data(mtcars)
X < -as.matrix(mtcars[, -1])
Y <- mtcars$mpg
X \leq cbind(1, X)
# Normal Equation
beta <- solve(t(X) %*% X) %*% t(X) %*% Y
cat("Regression Coefficients:\n"); print(beta)
```

```
# PCR
per model <- per(mpg ~ ., data = mtcars, validation = "CV")
summary(pcr model)
pred <- predict(pcr model, ncomp = 5)</pre>
cat("RMSE:", sqrt(mean((Y - pred)^2)), "\n")
13.Explore conditional probabilities, independence, and expected outliers. Examine
distribution shapes, skewness, correlations
library(ggplot2); library(corrplot); library(e1071)
data(mtcars)
# Conditional Probability
pAandB <- sum(mtcars$mpg > 20 & mtcars$hp > 100) / nrow(mtcars)
pB <- sum(mtcars$hp > 100) / nrow(mtcars)
cat("P(High MPG | HP > 100):", pAandB / pB, "\n")
# Independence (Correlation)
cat("Correlation (wt vs hp):", cor(mtcars$wt, mtcars$hp), "\n")
# Outliers using Z-score
z <- scale(mtcars$mpg)
cat("Outliers:", which(abs(z) > 3), "\n")
# Skewness
cat("Skewness of MPG:", skewness(mtcars$mpg), "\n")
# Histogram
ggplot(mtcars, aes(mpg)) + geom histogram(binwidth = 2, fill = "skyblue", color = "black")
# Correlation Heatmap
corrplot(cor(mtcars), method = "circle")
```

14. calculate the regression coefficients to fit the model. Discuss how multicollinearity among predictor variables affects the reliability

```
data(mtcars)
X \le -as.matrix(mtcars[, -1])
Y <- mtcars$mpg
X \leq cbind(1, X)
# Normal Equation
beta <- solve(t(X) %*% X) %*% t(X) %*% Y
cat("Regression Coefficients:\n"); print(beta)
15. Show the relationship between cyl (Cylinders) and hp (Horsepower) using a
box/scatter plot.
library(ggplot2)
data(mtcars)
mtcars$cyl <- as.factor(mtcars$cyl)</pre>
# Box Plot
ggplot(mtcars, aes(x = cyl, y = hp)) +
 geom boxplot(fill = "skyblue") +
 labs(title = "Box Plot: Cylinders vs Horsepower")
# Scatter Plot
ggplot(mtcars, aes(x = cyl, y = hp)) +
 geom jitter(width = 0.1, color = "blue") +
 labs(title = "Scatter Plot: Cylinders vs Horsepower")
16. Visualize correlations, split data, and train a Logistic Regression model. Evaluate
using accuracy, precision, recall, and F1-score.
library(ggplot2)
library(caret)
library(ggcorrplot)
```

```
library(pROC)
data(mtcars)
# Correlation Heatmap
cor matrix <- cor(mtcars)
ggcorrplot(cor matrix, hc.order = TRUE, type = "lower", lab = TRUE)
# Create binary target
mtcars$mpg binary <- ifelse(mtcars$mpg > 20, 1, 0)
# Train-test split
set.seed(123)
trainIndex <- createDataPartition(mtcars$mpg binary, p = 0.8, list = FALSE)
train data <- mtcars[trainIndex, ]</pre>
test data <- mtcars[-trainIndex, ]
# Logistic Regression
model <- glm(mpg binary ~ ., data = train data, family = "binomial")
probs <- predict(model, test data, type = "response")</pre>
pred < -ifelse(probs > 0.5, 1, 0)
# Evaluation
cm <- confusionMatrix(factor(pred), factor(test data$mpg binary))
print(cm)
cat("Accuracy: ", cm$overall['Accuracy'], "\n")
cat("Precision: ", cm$byClass['Pos Pred Value'], "\n")
cat("Recall: ", cm$byClass['Sensitivity'], "\n")
cat("F1 Score: ", 2 * ((cm$byClass['Pos Pred Value'] * cm$byClass['Sensitivity']) /
```

```
(cm$byClass['Pos Pred Value'] + cm$byClass['Sensitivity'])), "\n")
```

17.Identify numerical/categorical variables and handle missing values by imputing with mean, median, or mode. Convert categorical data into factors

```
library(dplyr)
data(mtcars)
# Introduce missing value
mtcars[1, "mpg"] <- NA
# Impute numeric with mean
mtcars$mpg[is.na(mtcars$mpg)] <- mean(mtcars$mpg, na.rm = TRUE)
# Mode function
get mode <- function(x) {
 ux <- unique(x)
 ux[which.max(tabulate(match(x, ux)))]
}
# Convert & Impute categorical with mode
mtcars$cyl <- as.factor(mtcars$cyl)</pre>
mtcars$gear <- as.factor(mtcars$gear)
mtcars$carb <- as.factor(mtcars$carb)</pre>
# Print structure
str(mtcars)
18 .Visualize mpg vs. wt (Weight) with color for cyl. 5.Create a bar plot showing car
counts by gear (Transmission gears).
library(ggplot2)
data(mtcars)
```

```
# Scatter plot
ggplot(mtcars, aes(x = wt, y = mpg, color = as.factor(cyl))) +
 geom point(size = 3) +
 labs(title = "mpg vs wt", color = "Cylinders") +
 theme minimal()
# Bar plot
ggplot(mtcars, aes(x = as.factor(gear))) +
 geom bar(fill = "skyblue", color = "black") +
 labs(title = "Car Counts by Gear", x = "Gears", y = "Count") +
 theme minimal()
19. handle missing values by imputing with mean, median, or mode.
library(dplyr)
data(mtcars)
# Introduce missing values
mtcars[1, "mpg"] <- NA
mtcars[3, "cyl"] <- NA
# Impute numeric
mtcars$mpg[is.na(mtcars$mpg)] <- mean(mtcars$mpg, na.rm = TRUE)
mtcars$cyl[is.na(mtcars$cyl)] <- median(mtcars$cyl, na.rm = TRUE)
# Mode function
get mode <- function(x) {
 ux <- unique(x)
 ux[which.max(tabulate(match(x, ux)))]
}
```

```
# Impute categorical
mtcars$gear[is.na(mtcars$gear)] <- get_mode(mtcars$gear)
# Show final data
print(mtcars)
20. Apply PCA for dimensionality reduction and assess if performance improves for
better
library(caret)
library(e1071)
data(mtcars)
# Standardize data
scaled data <- scale(mtcars)</pre>
# PCA
pca <- prcomp(scaled data, center = TRUE, scale. = TRUE)
summary(pca)
screeplot(pca, main = "Scree Plot", col = "blue")
# Keep first 2 components
pca data <- as.data.frame(pca$x[, 1:2])
pca data$mpg <- mtcars$mpg
# Train/test split
set.seed(123)
index <- createDataPartition(pca_data$mpg, p = 0.8, list = FALSE)
train <- pca data[index,]
test <- pca_data[-index, ]
```

```
# Model on PCA data

model <- glm(mpg ~ ., data = train, family = "gaussian")

pred <- predict(model, test)

rmse <- sqrt(mean((pred - test$mpg)^2))

print(paste("RMSE (PCA model):", rmse))
```